

Carl Glahn

SID: 998092557

ECS 171

### Report for HW 3

Notes for code:

I put this in the code as well but credit for all of the m files in the MSVM pack are not mine however they are public and the authors asked to be cited thusly:

*F. Lauer and Y. Guermeur, MSVMpack: a Multi-Class Support Vector Machine Package, Journal of Machine Learning Research, 12:2269-2272, 2011.*

To run the code you might need to change the `addpath()` function although it should work as written with the `"addpath(genpath('MSVMpack1.5/matlab'))"`

Also I had to do some work arounds to get the ROC and PR curves with the toolbox I chose, the details of which are further explained in #4

1. For this problem I decided to use ridge regression, which penalizes the objective function (RSS) by adding the sum of the weights squared multiplied by a weight constrained parameter ( $\lambda$ ) while also bounding the sum of the squared weights to be below a constant  $c$  which I chose to be one. What this does is that it creates a radius in which the weights are allowed to take values and the features that have the most effect on the outputs will take most of total value of the sum of squared weights and the other feature's weights will drop to around zero.

To find my optimal  $\lambda$  I set it to a predefined value (I started at one), generated my weights, and then use the error to adjust  $\lambda$ . What I found was that  $\lambda$  settled at a value of .15, where the error would bounce from positive and negative  $7.0897e-06$ . To save time on the program I preset  $\lambda$  to .15, although my program would alter  $\lambda$  if the error was big enough (although after printing out some values I found this never occurred).

Because the shape of ridge regression is circular, many of the weights never quite drop to 0 and instead settle around .001. To fix this I set a threshold where if the weight was below .001, it would be treated as if it the weight was a zero. Given this the (lowest) number of non-zero features (NNZ) in the cross validation was 2329.

The table for the 10 fold cross validation errors are below and the generalization error (the average) was .024994

Error	0.012782	0.049608	0.014803	0.043084	0.003046	0.025071	0.01218	0.003081	0.068619	0.017663
NNZ elements	2505	2614	2553	2619	2629	2592	2453	2566	2329	2553

2. For this problem I took 100 random samplings of the original dataset (bootstraps) and ran them through my cross validation linear regression algorithm. While doing this I stored all of the weights that were generated for each bootstrap data set. After I had generate rated 100 sets of weights I then ran a sample of the original dataset through every set of weights and store the outputted predictions in a vector. I would then take the mean of the vector and use it to generate the upper and lower bounds (for a 95% confidence interval) for that sample. I would then repeat this process for each sample of the dataset. To do this I had to assume that the values of the predictions generated by the set of weights of the bootstrap dataset had a normal distribution. The results are below

Sample Number	Lower Bound	Upper Bound
1	0.623955276	0.629381937
2	0.636971178	0.64168564
3	0.628498858	0.632744783
4	0.725743885	0.728205221
5	0.715667235	0.71966343
6	0.723075155	0.725367968
7	0.465756879	0.467394772
8	0.467521634	0.469468582
9	0.465532523	0.468148572
10	0.477974076	0.479421874
11	0.481011655	0.482419862
12	0.477667248	0.479668891
13	0.470306588	0.477712149
14	0.470306588	0.477712149
15	0.461478801	0.483432247
16	0.462230949	0.479513541
17	0.479958529	0.490165307
18	0.464428968	0.479674003
19	0.470899709	0.483029836
20	0.47515208	0.486222081
21	0.380244577	0.432207253
22	0.681166107	0.791083669
23	0.491731483	0.499281829
24	0.314777571	0.40524363
25	0.381527591	0.387037198
26	0.376139749	0.381199142
27	0.381938634	0.388340175
28	0.406091451	0.415754596
29	0.416013544	0.423903978
30	0.420306464	0.429216426

31	0.477818991	0.480758979
32	0.475866602	0.479116793
33	0.479274017	0.484096546
34	0.487570637	0.490538935
35	0.489471022	0.49445035
36	0.480557024	0.483592335
37	0.63155082	0.634559284
38	0.623774869	0.629578995
39	0.63095054	0.633855536
40	0.635499167	0.640162751
41	0.634655283	0.641379787
42	0.587531729	0.605470481
43	0.678004018	0.682241779
44	0.647292894	0.66881909
45	0.715262552	0.720814335
46	0.708279596	0.711745101
47	0.688873092	0.697052302
48	0.2190242	0.262315114
49	0.211931577	0.275076856
50	0.296984594	0.306159353
51	0.252727132	0.272825354
52	0.444332975	0.460132579
53	0.440594997	0.456560808
54	0.444635472	0.46131852
55	0.45034981	0.536307209
56	0.1118334	0.125807637
57	0.112826833	0.126884945
58	0.243258286	0.249001839
59	0.354575121	0.359800445
60	0.231637856	0.236726769
61	0.233449869	0.238727263
62	0.23125939	0.236416996
63	0.265059869	0.30562457
64	0.253544477	0.29273402
65	0.239069469	0.271184966
66	0.218778887	0.245718868
67	0.374779148	0.38173773
68	0.371136309	0.377050458
69	0.380740858	0.391625325
70	0.369513055	0.372433518
71	0.361576299	0.365623804

72	0.366221606	0.369587031
73	0.373691002	0.377128749
74	0.368716902	0.371021683
75	0.369641223	0.372292366
76	0.210759875	0.21965342
77	0.239704451	0.247966025
78	0.231496912	0.23739093
79	0.221710993	0.227780187
80	0.24102472	0.24954876
81	0.246792697	0.258681361
82	0.120079633	0.137064388
83	0.131592288	0.145463073
84	0.157610525	0.161361202
85	0.095454221	0.116910744
86	0.129296772	0.138701447
87	0.100876305	0.116939673
88	0.117813094	0.127180649
89	0.109397834	0.115045047
90	0.085554879	0.097420064
91	0.090377759	0.10025225
92	0.096246111	0.110590907
93	0.083155836	0.089873902
94	0.087950965	0.101810994
95	1.215212277	1.236977803
96	1.157085236	1.18945105
97	1.160223625	1.201299294
98	1.215751999	1.230546365
99	1.23914671	1.262265263
100	0.64335863	0.649968401
101	0.619308159	0.63437775
102	0.634579485	0.645957108
103	0.647742925	0.655235096
104	0.617699643	0.634967292
105	0.640755197	0.649841165
106	0.240817178	0.266748358
107	0.218186046	0.254739047
108	0.272672534	0.289193668
109	0.151343548	0.162326111
110	0.176853198	0.193447736
111	0.182782195	0.202544556
112	0.351274982	0.417796756


113	0.317938463	0.350240242
114	0.292816976	0.301218806
115	0.182306049	0.229557043
116	0.068157468	0.027389205
117	0.162212971	0.171588564
118	0.107648726	0.112309823
119	0.113333898	0.124783609
120	0.135874222	0.171196813
121	0.104311494	0.107889757
122	0.105917518	0.110522682
123	0.187449019	0.246195696
124	0.104368334	0.107665896
125	0.106020467	0.109795215
126	0.10370182	0.106892606
127	0.103133055	0.106590846
128	0.103707606	0.106719067
129	0.158811266	0.201600121
130	0.534491859	0.567119925
131	0.632833241	0.65932389
132	0.613046119	0.640118763
133	0.54855408	0.573234234
134	0.582699816	0.600321666
135	0.596987731	0.618173833
136	0.507854456	0.550246986
137	0.611013174	0.648124816
138	0.538703576	0.563752818
139	0.545889215	0.569846065
140	0.135946901	0.165844446
141	0.571156447	0.584732957
142	0.727102716	0.769982316
143	0.408004279	0.414042333
144	0.400126204	0.402088922
145	0.398024959	0.399356561
146	0.398534517	0.400018407
147	0.394412007	0.396662323
148	0.400949959	0.403632165
149	0.397727181	0.398891167
150	0.394292244	0.397256276
151	0.401295867	0.402511679
152	0.40000133	0.40119648
153	0.39601449	0.397791979

154	0.400099999	0.40131648
155	0.398107438	0.400864515
156	0.392382409	0.396570238
157	0.408255173	0.414504982
158	0.098002824	0.135376061
159	0.646192567	0.670142118
160	0.584676244	0.599949722
161	0.588647453	0.602649949
162	0.29441013	0.32282675
163	0.314142594	0.343057924
164	0.338454977	0.377563963
165	0.28734361	0.317416691
166	0.172445885	0.242443613
167	0.288728871	0.312423502
168	0.225920284	0.30654822
169	0.335029002	0.371738735
170	0.324746942	0.344119793
171	0.31753668	0.340935633
172	0.557247555	0.578039983
173	0.337648096	0.369646298
174	0.294534286	0.310532221
175	0.561286448	0.5783625
176	0.537100752	0.568001016
177	0.516186094	0.552625651
178	0.514065573	0.52834214
179	0.459054825	0.477763971
180	0.197052103	0.198813439
181	0.200173278	0.201165134
182	0.200820248	0.202305287
183	0.197230249	0.203833356
184	0.200737671	0.207506921
185	0.197591813	0.200149403
186	0.199263676	0.200618192
187	0.167048593	0.180510259
188	0.201529499	0.203328722
189	0.196887699	0.205843311
190	0.19585881	0.206644513
191	0.192251777	0.196566788
192	0.202226646	0.204326338
193	0.199392958	0.200220997
194	0.198192502	0.199553888


3. For this problem I trained on the whole dataset to get the most accurate result. After generating the weights and taking the mean of every column of the input, the predicted growth rate outputted was .3936.
4. For this problem I used 4 multi-classification SVMs (one for each classification categories). To eliminate features I trained using ridge regression on the whole data set and picked the features that had a weight coefficient of greater than .001 of which there were 1561. I graphed the ROC curves in the preferred method using 10 points for confidence thresholds of .5, .55,...,.95 as recommended. However there was one change and one adjustment I made to this method. The change was that there was one point for the test run for each threshold, meaning that I took the generated labels for a particular threshold (after the whole cross validation was done) and calculated one value for FPR, TPR, recall, and precision. The adjustment I had to make came from the fact that the toolbox I was using did not give me the option of adjusting a threshold. However it did give me the confidence ratings of each class of prediction by outputting Mathew's coefficients (one set for each fold of the cross validation) which essentially measure correlation on a scale from -1 to 1. Using these coefficients I could mimic having thresholds by referring to the coefficients to see how much of a correlation there was for that class. I set the initial benchmark for a .5 threshold at a .5 (half) correlation. This process was further complicated because the toolbox only outputted the coefficients to the console and not in a vector, meaning that I was forced to manually copy the coefficients over to a text file and then read them in from there. This being said if you just run my code without manually copying the coefficients from the console to the text file (they change every time because of the random ordering) you will get a very high error because the old coefficients are no longer correct. To get similar results as are reported below you must run the SVMs while commenting out the graphing function, then copy over the coefficients manually, then uncomment the graphing function and run the code again commenting out the SVMs . Obviously this isn't the cleanest way to do things but it did give me the required results. Which are displayed below.

	Strain	Medium	Environment	Gene Perturbation
AUC	0.0208	0.0607	0.0452	0.0245
AUPRC	0.3941	0.195	0.5428	0.5042


Strain Curves: (Error of .1186)

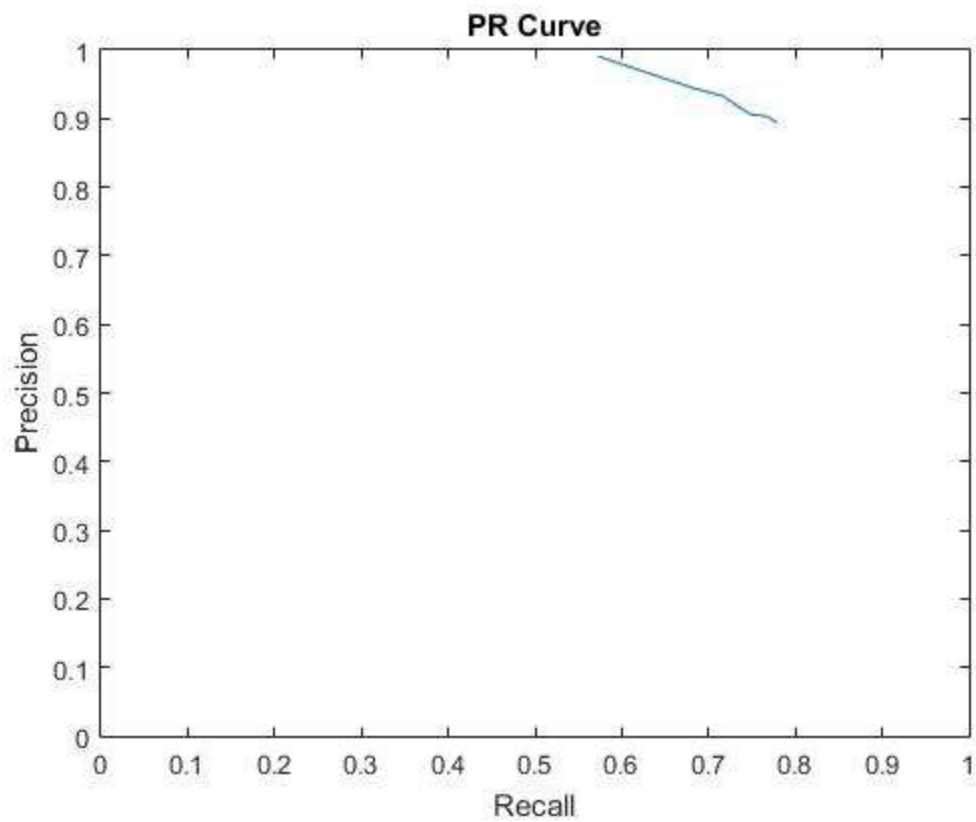







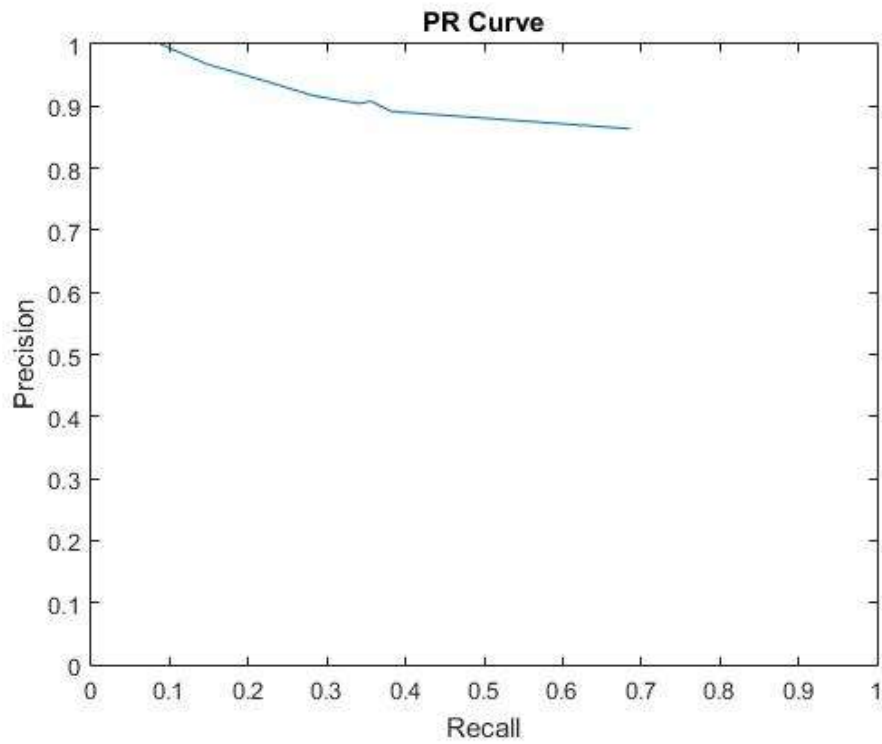
Medium Curves: (Error of .1959)







Environment Curves: (Error of .2062)





Gene Perturbation Curves: (Error of .1804)





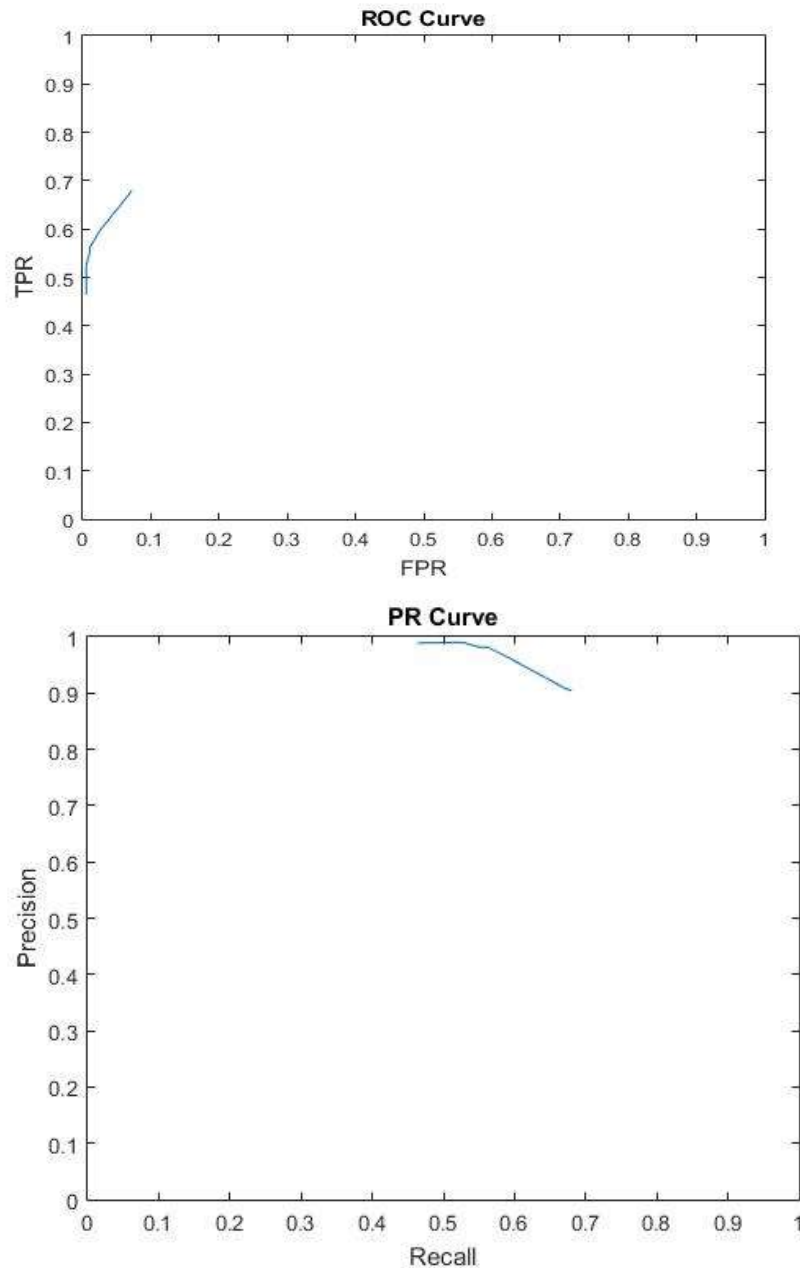
One thing I did notice was that my AUC and AUPRC values seemed very low considering that, for the most part, I had achieved an (approximate) 80% accuracy rating on each SVM.

5. For this problem I found that there was a total of 23 different pairing between the environmental and medium labeling. Naturally I simply encoded these labeling from 1 to 23 and graphed things in the same manner as I did for problem 4.

Error: .2474

AUC: .0414

AUPRC: .2085



As you can see not only did the composite predictor have a higher error than the 2 separate predictors, it also had a lower AUC value. From this we can conclude that we are in fact better off with 2 predictors rather than one composite one. There were 23 classes in total for the composite predictor so the baseline accuracy would be .04347 ( $1/23$ ). As before, I was surprised by the low AUC and AUPRC values considering the (approximate) 75% accuracy rating.


6. For this question I did the PCA analysis for the composite SVM as the instructions said to do this for "the SVM classifier". The TA later mentioned that if one had already done this problem in this


manner, that we could keep our results as long as we put a note for it in the report. The results are displayed below.

Error: .7938

AUC: .0111

AUPRC: .0211





After looking at the statistics and figures above it is painfully obvious that the three principle components did not retain nearly enough information about the original 4495 features to produce accurate results (the error margin alone increased be 50%).